

# DSC 441 – FUNDAMENTAL OF DATA SCIENCE

## HOMEWORK 4

NAME – Goutham Selvakumar

### PROBLEM 1

For this problem, you will tune and apply kNN and compare it to other classifiers. We will use the wine quality data, which has a number of measurements about chemical components in wine, plus a quality rating. There are separate files for red and white wines, so the first step is some data preparation

- a. For this problem, you will tune and apply kNN and compare it to other classifiers. We will use the wine quality data, which has a number of measurements about chemical components in wine, plus a quality rating. There are separate files for red and white wines, so the first step is some data preparation.

```
```{r}
# Importing the datasets and separating them with a semicolon
winequality_white <- read.csv("winequality-white.csv", sep = ";")
winequality_red <- read.csv("winequality-red.csv", sep = ";")
summary(winequality_white)
```
```

| fixed.acidity  | volatile.acidity | citric.acid     | residual.sugar | chlorides        |
|----------------|------------------|-----------------|----------------|------------------|
| Min. : 3.800   | Min. : 0.0800    | Min. : 0.0000   | Min. : 0.600   | Min. : 0.00900   |
| 1st Qu.: 6.300 | 1st Qu.: 0.2100  | 1st Qu.: 0.2700 | 1st Qu.: 1.700 | 1st Qu.: 0.03600 |
| Median : 6.800 | Median : 0.2600  | Median : 0.3200 | Median : 5.200 | Median : 0.04300 |
| Mean : 6.855   | Mean : 0.2782    | Mean : 0.3342   | Mean : 6.391   | Mean : 0.04577   |
| 3rd Qu.: 7.300 | 3rd Qu.: 0.3200  | 3rd Qu.: 0.3900 | 3rd Qu.: 9.900 | 3rd Qu.: 0.05000 |
| Max. : 14.200  | Max. : 1.1000    | Max. : 1.6600   | Max. : 65.800  | Max. : 0.34600   |

| free.sulfur.dioxide | total.sulfur.dioxide | density         | pH             | sulphates       |
|---------------------|----------------------|-----------------|----------------|-----------------|
| Min. : 2.00         | Min. : 9.0           | Min. : 0.9871   | Min. : 2.720   | Min. : 0.2200   |
| 1st Qu.: 23.00      | 1st Qu.: 108.0       | 1st Qu.: 0.9917 | 1st Qu.: 3.090 | 1st Qu.: 0.4100 |
| Median : 34.00      | Median : 134.0       | Median : 0.9937 | Median : 3.180 | Median : 0.4700 |
| Mean : 35.31        | Mean : 138.4         | Mean : 0.9940   | Mean : 3.188   | Mean : 0.4898   |
| 3rd Qu.: 46.00      | 3rd Qu.: 167.0       | 3rd Qu.: 0.9961 | 3rd Qu.: 3.280 | 3rd Qu.: 0.5500 |
| Max. : 289.00       | Max. : 440.0         | Max. : 1.0390   | Max. : 3.820   | Max. : 1.0800   |

| alcohol        | quality        |
|----------------|----------------|
| Min. : 8.00    | Min. : 3.000   |
| 1st Qu.: 9.50  | 1st Qu.: 5.000 |
| Median : 10.40 | Median : 6.000 |
| Mean : 10.51   | Mean : 5.878   |
| 3rd Qu.: 11.40 | 3rd Qu.: 6.000 |
| Max. : 14.20   | Max. : 9.000   |

```
summary(winequality_red)
...

fixed.acidity    volatile.acidity    citric.acid    residual.sugar    chlorides
Min.   : 4.60    Min.   :0.1200    Min.   :0.000    Min.   : 0.900    Min.   :0.01200
1st Qu.: 7.10    1st Qu.:0.3900    1st Qu.:0.090    1st Qu.: 1.900    1st Qu.:0.07000
Median : 7.90    Median :0.5200    Median :0.260    Median : 2.200    Median :0.07900
Mean   : 8.32    Mean   :0.5278    Mean   :0.271    Mean   : 2.539    Mean   :0.08747
3rd Qu.: 9.20    3rd Qu.:0.6400    3rd Qu.:0.420    3rd Qu.: 2.600    3rd Qu.:0.09000
Max.   :15.90    Max.   :1.5800    Max.   :1.000    Max.   :15.500    Max.   :0.61100

free.sulfur.dioxide total.sulfur.dioxide    density    pH    sulphates
Min.   : 1.00    Min.   : 6.00    Min.   :0.9901    Min.   :2.740    Min.   :0.3300
1st Qu.: 7.00    1st Qu.: 22.00    1st Qu.:0.9956    1st Qu.:3.210    1st Qu.:0.5500
Median :14.00    Median : 38.00    Median :0.9968    Median :3.310    Median :0.6200
Mean   :15.87    Mean   :46.47    Mean   :0.9967    Mean   :3.311    Mean   :0.6581
3rd Qu.:21.00    3rd Qu.: 62.00    3rd Qu.:0.9978    3rd Qu.:3.400    3rd Qu.:0.7300
Max.   :72.00    Max.   :289.00    Max.   :1.0037    Max.   :4.010    Max.   :2.0000

alcohol    quality
Min.   : 8.40    Min.   :3.000
1st Qu.: 9.50    1st Qu.:5.000
Median :10.20    Median :6.000
Mean   :10.42    Mean   :5.636
3rd Qu.:11.10    3rd Qu.:6.000
Max.   :14.90    Max.   :8.000
```

Now, check the type of data in the columns of the wines dataset

```
> typeof(winequality_red$chlorides)
[1] "double"
> typeof(winequality_red$citric.acid)
[1] "double"
> typeof(winequality_red$residual.sugar)
[1] "double"
> typeof(winequality_red$alcohol)
[1] "double"
```

Note that there is no NAs in the dataset. Now, let's add a column for the types of wines.

```
# Add the column for type of wine
winequality_red$type <- c('red')
winequality_white$type <- c('white')
# Now combine both tables using the full_join
wines <- full_join(winequality_red, winequality_white)
head(wines)
```


Description: df [6 x 13]

|   | fixed.acidity<br><dbl> | volatile.acidity<br><dbl> | citric.acid<br><dbl> | residual.sugar<br><dbl> | chlorides<br><dbl> | free.sulfur.dioxide<br><dbl> | total.sulfur.dioxide<br><dbl> | density<br><dbl> | pH<br><dbl> |
|---|------------------------|---------------------------|----------------------|-------------------------|--------------------|------------------------------|-------------------------------|------------------|-------------|
| 1 | 7.4                    | 0.70                      | 0.00                 | 1.9                     | 0.076              | 11                           | 34                            | 0.9978           | 3.51        |
| 2 | 7.8                    | 0.88                      | 0.00                 | 2.6                     | 0.098              | 25                           | 67                            | 0.9968           | 3.20        |
| 3 | 7.8                    | 0.76                      | 0.04                 | 2.3                     | 0.092              | 15                           | 54                            | 0.9970           | 3.26        |
| 4 | 11.2                   | 0.28                      | 0.56                 | 1.9                     | 0.075              | 17                           | 60                            | 0.9980           | 3.16        |
| 5 | 7.4                    | 0.70                      | 0.00                 | 1.9                     | 0.076              | 11                           | 34                            | 0.9978           | 3.51        |
| 6 | 7.4                    | 0.66                      | 0.00                 | 1.8                     | 0.075              | 13                           | 40                            | 0.9978           | 3.51        |

6 rows | 1-10 of 13 columns

- b. Use PCA to create a projection of the data to 2D and show a scatterplot with color showing the wine type.**

```
```{r}
# Create the Dummies for the dataset
dummy <- dummyVars(type ~ ., data = wines)
dummies <- as.data.frame(predict(dummy, newdata = wines))
set.seed(123)
# Calculate PCA
pca <- prcomp(dummies)
# Save as data frame
rotated_data <- as.data.frame(pca$x)
# Add the original label 'type' as a reference
rotated_data$Color <- wines$type
# Plot and color the labels based on wine type red (or) white
ggplot(data = rotated_data, aes(x = PC1, y = PC2, col = Color)) + geom_point(alpha = 0.3)
```
```



- c. We are going to try kNN, SVM and decision trees on this data. Based on the ‘shape’ of the data in the visualization from (b), which do you think will do best and why?

I think the best method in this case may be to use KNN. The wines dataset is somewhat large so KNN may be faster and more efficient solution. Also, the non-linear relationship will allow for a more accurate prediction by using the nearest K neighbour to predict the type of wine. I’ve performed some sample work below that helped me also get to this conclusion, although SVM and KNN performed similarly in regards to accuracy. However, SVM’s patterns may be limited in this case (we can discover this later).

#### Finding the K nearest neighbour:

```
```{r}
set.seed(123)
# Scaling is crucial for KNN
ctrl <- trainControl(method = "cv", number = 10)
knnFit <- train(type ~ ., data = wines,
               method = "knn",
               trControl = ctrl,
               preProcess = c("center", "scale"))
# Output for KNN fit
knnFit
```
```

```

k-Nearest Neighbors

6497 samples
 12 predictor
 2 classes: 'red', 'white'

Pre-processing: centered (12), scaled (12)
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 5848, 5847, 5847, 5848, 5847, 5848, ...
Resampling results across tuning parameters:

k  Accuracy  Kappa
5  0.9923051  0.9792070
7  0.9929205  0.9808827
9  0.9930741  0.9813154

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was k = 9.

```

## **SVM:**

```

# Fit the Model
svm1 <- train(type ~., data = wines, method = "svmLinear")
# Evaluate the Fit
svm1

```

```

Support Vector Machines with Linear Kernel

6497 samples
 12 predictor
 2 classes: 'red', 'white'

No pre-processing
Resampling: Bootstrapped (25 reps)
Summary of sample sizes: 6497, 6497, 6497, 6497, 6497, 6497, ...
Resampling results:

Accuracy  Kappa
0.994937  0.9862915

Tuning parameter 'C' was held constant at a value of 1

```

## **Decision Tree:**

```

# Evaluation Method
train_control = trainControl(method = "cv", number = 10)
# Fit the Model
tree1 <- train(type ~., data = wines, method = "rpart", trControl = train_control)
tree1

```

```

CART

6497 samples
 12 predictor
 2 classes: 'red', 'white'

No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 5847, 5848, 5847, 5847, 5848, 5847, ...
Resampling results across tuning parameters:

cp          Accuracy  Kappa
0.06253909  0.9465933  0.8498495
0.06754221  0.9348963  0.8147365
0.70043777  0.7855788  0.1500825

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was cp = 0.06253909.

```

- d. Use kNN (tune k), use decision trees (basic rpart method is fine), and SVM (tune C) to predict type from the rest of the variables. Compare the accuracy values – is this what you expected? Can you explain it? Note: you will need to fix the columns names for rpart because it is not able to handle the underscores. This code will do the trick (assuming you called your data wine\_quality):
- ```
colnames(wine_quality) <- make.names(colnames(wine_quality)).
```

### KNN:

```
```{r}
set.seed(123)
ctrl <- trainControl(method="cv", number = 10)
knnFit <- train(type ~., data = wines,
                method = "knn",
                trControl = ctrl,
                preProcess = c("center", "scale"),
                tuneLength = 15)

knnFit
```
```

k-Nearest Neighbors

6497 samples  
12 predictor  
2 classes: 'red', 'white'

Pre-processing: centered (12), scaled (12)  
Resampling: Cross-Validated (10 fold)  
Summary of sample sizes: 5847, 5847, 5848, 5849, 5847, 5847, ...  
Resampling results across tuning parameters:

| k  | Accuracy  | Kappa     |
|----|-----------|-----------|
| 5  | 0.9923048 | 0.9792105 |
| 7  | 0.9936890 | 0.9829507 |
| 9  | 0.9933813 | 0.9821286 |
| 11 | 0.9927659 | 0.9805034 |
| 13 | 0.9926121 | 0.9801038 |
| 15 | 0.9923044 | 0.9792783 |
| 17 | 0.9927659 | 0.9804950 |
| 19 | 0.9926116 | 0.9800879 |
| 21 | 0.9929198 | 0.9809104 |
| 23 | 0.9924582 | 0.9796711 |
| 25 | 0.9923039 | 0.9792570 |
| 27 | 0.9923044 | 0.9792539 |
| 29 | 0.9923044 | 0.9792539 |
| 31 | 0.9916885 | 0.9776041 |
| 33 | 0.9918428 | 0.9780181 |

Accuracy was used to select the optimal model using the largest value.  
The final value used for the model was k = 7.

### For k = 7:

```
fit <- kmeans(dummies, centers = 7, nstart = 25)
# Display the kmeans object information
fit
```

K-means clustering with 7 clusters of sizes 856, 910, 354, 717, 1382, 1131, 1147

```
Cluster means:
fixed,acidity volatile,acidity citric,acid residual,sugar chlorides free,sulfur,dioxide
1 6.939486 0.286596 0.351612 5.869974 0.05149224 48.915904
2 9.448871 0.512174 0.271495 2.450200 0.0265165 9.10989
3 7.030791 0.306963 0.3606215 10.179802 0.05212147 55.213277
4 7.723570 0.4324407 0.2865930 2.613595 0.07973640 19.561590
5 6.836459 0.237152 0.324765 5.246253 0.04537192 30.159198
6 6.924637 0.314893 0.3086207 3.793143 0.04803006 24.540472
7 6.920837 0.282125 0.3409765 7.195074 0.04927114 39.075414
total,sulfur,dioxide density pH sulphates alcohol quality
1 182.81600 0.9960240 3.188166 0.8109229 9.790479 5.896449
2 25.89022 0.9964969 3.303615 0.6385824 10.592271 5.718681
3 224.39972 0.9967192 3.173107 0.5207345 9.545763 5.325424
4 59.94421 0.9949663 3.269243 0.6056206 10.662366 5.647141
5 122.64146 0.9934417 3.195166 0.4963749 10.770466 6.013025
6 95.01551 0.9928337 3.209054 0.5002019 11.033687 5.979664
7 152.05929 0.9947757 3.139355 0.4900697 10.250753 5.874455
```

```
Clustering vector:
1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19
2 4 4 4 2 2 2 2 2 2 6 4 6 4 2 7 7 6 4 2
20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38
4 4 4 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57
2 6 6 4 2 2 2 4 5 2 2 6 2 2 2 5 4 4 2
58 59 60 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76
5 4 4 4 5 2 4 2 2 2 2 4 2 2 6 6 4 6 2
77 78 79 80 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95
7 2 2 2 2 4 4 4 2 5 2 5 2 5 2 7 5 5 2
96 97 98 99 100 101 102 103 104 105 106 107 108 109 110 111 112 113 114
6 2 2 2 4 4 6 4 2 4 6 4 7 6 2 4 2 4 2
115 116 117 118 119 120 121 122 123 124 125 126 127 128 129 130 131 132 133
4 4 2 2 2 4 6 4 2 2 2 6 6 2 2 2 2 5 6 6
134 135 136 137 138 139 140 141 142 143 144 145 146 147 148 149 150 151 152
2 2 4 5 4 6 6 4 4 4 4 4 7 6 6 2 4 2 4
153 154 155 156 157 158 159 160 161 162 163 164 165 166 167 168 169 170 171
6 6 5 5 5 5 2 6 2 2 2 5 5 6 6 4 4 4 2
172 173 174 175 176 177 178 179 180 181 182 183 184 185 186 187 188 189 190
2 2 4 2 4 2 2 2 2 2 2 5 4 4 4 6 4 2 5 7
191 192 193 194 195 196 197 198 199 200 201 202 203 204 205 206 207 208 209
5 4 5 2 4 2 4 2 4 2 2 2 2 2 2 2 2 2 5 6
210 211 212 213 214 215 216 217 218 219 220 221 222 223 224 225 226 227 228
2 2 4 4 4 2 2 5 2 2 4 7 4 4 2 2 4 4 2 2
229 230 231 232 233 234 235 236 237 238 239 240 241 242 243 244 245 246 247
4 2 6 4 4 2 2 2 2 2 2 2 2 2 2 2 2 2 2
248 249 250 251 252 253 254 255 256 257 258 259 260 261 262 263 264 265 266
2 4 2 2 4 2 2 2 2 2 4 2 4 2 4 2 4 2 2
267 268 269 270 271 272 273 274 275 276 277 278 279 280 281 282 283 284 285
4 2 2 2 4 2 2 4 2 2 4 2 4 2 2 2 4 2 2
286 287 288 289 290 291 292 293 294 295 296 297 298 299 300 301 302 303 304
4 4 4 4 4 6 4 2 2 2 2 2 2 2 2 2 2 2 2
305 306 307 308 309 310 311 312 313 314 315 316 317 318 319 320 321 322 323
6 2 4 2 2 2 2 2 2 6 6 5 4 4 4 6 4 6 2
324 325 326 327 328 329 330 331 332 333 334 335 336 337 338 339 340 341 342
2 2 2 2 2 2 2 2 2 5 4 4 4 6 2 2 2 2
343 344 345 346 347 348 349 350 351 352 353 354 355 356 357 358 359 360 361
2 2 4 4 2 2 4 2 2 2 2 2 2 2 2 2 2 2 2
362 363 364 365 366 367 368 369 370 371 372 373 374 375 376 377 378 379 380
4 2 2 2 2 2 2 2 2 4 2 2 2 2 2 2 2 2 2
381 382 383 384 385 386 387 388 389 390 391 392 393 394 395 396 397 398 399
2 4 2 2 2 4 2 4 4 6 6 2 2 2 2 5 4 4 4
400 401 402 403 404 405 406 407 408 409 410 411 412 413 414 415 416 417 418
2 5 2 2 2 4 2 2 2 2 2 4 6 6 4 2 5 2 5
```

```
419 420 421 422 423 424 425 426 427 428 429 430 431 432 433 434 435 436 437
2 4 4 4 4 2 4 4 4 2 2 2 2 2 2 2 4 2 4
438 439 440 441 442 443 444 445 446 447 448 449 450 451 452 453 454 455 456
2 4 4 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
457 458 459 460 461 462 463 464 465 466 467 468 469 470 471 472 473 474 475
2 4 2 2 2 2 2 2 2 2 2 2 2 2 2 2 4 2 2
476 477 478 479 480 481 482 483 484 485 486 487 488 489 490 491 492 493 494
2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 4 2 2
495 496 497 498 499 500 501 502 503 504 505 506 507 508 509 510 511 512 513
6 2 2 5 2 2 2 2 6 6 6 2 2 2 2 4 4 4 4
514 515 516 517 518 519 520 521 522 523 524 525 526 527 528 529 530 531 532
2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
533 534 535 536 537 538 539 540 541 542 543 544 545 546 547 548 549 550 551
2 2 2 2 2 2 2 2 2 4 2 2 2 2 2 2 4 2 2
552 553 554 555 556 557 558 559 560 561 562 563 564 565 566 567 568 569 570
4 4 6 2 2 2 2 2 2 2 2 2 5 5 4 4 2 2 2 4
571 572 573 574 575 576 577 578 579 580 581 582 583 584 585 586 587 588 589
2 4 2 4 6 4 2 5 6 6 2 2 2 2 2 2 2 2 2 2
590 591 592 593 594 595 596 597 598 599 600 601 602 603 604 605 606 607 608
2 4 7 4 2 2 5 2 2 2 2 2 2 2 2 2 2 2 2
609 610 611 612 613 614 615 616 617 618 619 620 621 622 623 624 625 626 627
4 4 4 2 2 2 2 6 4 5 2 2 5 6 2 4 4 4 2
628 629 630 631 632 633 634 635 636 637 638 639 640 641 642 643 644 645 646
2 2 6 2 2 4 6 6 6 2 7 2 4 2 4 2 4 2 2
647 648 649 650 651 652 653 654 655 656 657 658 659 660 661 662 663 664 665
2 2 2 7 2 7 4 2 2 4 2 4 2 2 2 2 2 2 4
666 667 668 669 670 671 672 673 674 675 676 677 678 679 680 681 682 683 684
4 2 2 2 2 2 7 2 2 2 2 2 2 6 4 2 2 2 2
685 686 687 688 689 690 691 692 693 694 695 696 697 698 699 700 701 702 703
7 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
704 705 706 707 708 709 710 711 712 713 714 715 716 717 718 719 720 721 722
6 2 4 2 2 2 2 6 5 2 4 4 2 4 2 4 2 2 6
723 724 725 726 727 728 729 730 731 732 733 734 735 736 737 738 739 740 741
4 5 2 2 4 2 4 2 2 2 2 2 2 2 2 2 2 2 2
742 743 744 745 746 747 748 749 750 751 752 753 754 755 756 757 758 759 760
5 2 6 6 2 4 4 4 2 2 2 2 2 2 2 2 2 2 2
761 762 763 764 765 766 767 768 769 770 771 772 773 774 775 776 777 778 779
6 2 2 2 4 4 4 6 6 2 6 7 7 2 2 2 2 2 2
780 781 782 783 784 785 786 787 788 789 790 791 792 793 794 795 796 797 798
6 4 2 6 2 6 4 4 4 6 5 6 6 2 2 4 4 4
799 800 801 802 803 804 805 806 807 808 809 810 811 812 813 814 815 816 817
2 6 6 4 4 2 2 2 2 2 2 2 2 2 2 2 2 2 2
818 819 820 821 822 823 824 825 826 827 828 829 830 831 832 833 834 835 836
6 4 2 4 2 4 2 2 2 2 2 2 2 2 2 2 2 2 2
837 838 839 840 841 842 843 844 845 846 847 848 849 850 851 852 853 854 855
537 538 539 540 541 542 543 544 545 546 547 548 549 550 551 552 553 554 555
6 5 2 2 2 2 4 6 4 2 2 2 2 2 2 2 2 2 2
856 857 858 859 860 861 862 863 864 865 866 867 868 869 870 871 872 873 874
2 4 4 2 2 2 2 6 6 6 6 6 6 6 6 6 6 6 6
875 876 877 878 879 880 881 882 883 884 885 886 887 888 889 890 891 892 893
2 2 2 4 4 6 2 2 2 2 2 2 2 2 2 2 2 2 2
894 895 896 897 898 899 900 901 902 903 904 905 906 907 908 909 910 911 912
6 4 4 4 4 2 2 2 2 2 2 2 2 2 2 2 2 2 2
913 914 915 916 917 918 919 920 921 922 923 924 925 926 927 928 929 930 931
2 2 2 2 2 2 2 6 4 6 4 2 2 2 2 2 2 2 2
932 933 934 935 936 937 938 939 940 941 942 943 944 945 946 947 948 949 950
2 4 2 2 4 4 2 2 2 2 2 2 2 2 2 2 2 2 2
951 952 953 954 955 956 957 958 959 960 961 962 963 964 965 966 967 968 969
2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
970 971 972 973 974 975 976 977 978 979 980 981 982 983 984 985 986 987 988
2 2 2 2 2 2 2 4 4 5 5 5 5 5 5 5 5 5
989 990 991 992 993 994 995 996 997 998 999 1000
2 2 2 4 4 4 6 4 2 2 2 2 2
[ reached getOption("max.print") -- omitted 5497 entries ]
```

```
Within cluster sum of squares by cluster:
[1] 330867.7 108090.8 367269.3 147020.0 301721.7 206162.7 340030.8
(between_SS / total_SS = 92.2 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss" "betweenss"
[7] "size"         "iter"         "ifault"
```

## SVM:

```
# I decided to use the Grid Search here to try different values of C
grid <- expand.grid(C = 10^seq(-5,2,0.5))
# Fit the Model
svm_grid <- train(type ~., data = wines, method = "svmLinear",
                  trControl = train_control, tuneGrid = grid)
# View grid search result
svm_grid
```

Support Vector Machines with Linear Kernel

```
6497 samples
12 predictor
2 classes: 'red', 'white'
```

```
No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 5847, 5847, 5847, 5847, 5847, ...
Resampling results across tuning parameters:
```

| C            | Accuracy  | Kappa       |
|--------------|-----------|-------------|
| 1.000000e-05 | 0.7538866 | 0.000000000 |
| 3.162278e-05 | 0.7541943 | 0.001878654 |
| 1.000000e-04 | 0.9302726 | 0.792017114 |
| 3.162278e-04 | 0.9839943 | 0.956278082 |
| 1.000000e-03 | 0.9906118 | 0.974623044 |
| 3.162278e-03 | 0.9915354 | 0.977177217 |
| 1.000000e-02 | 0.9926128 | 0.980992486 |
| 3.162278e-02 | 0.9938435 | 0.983410595 |
| 1.000000e-01 | 0.9946132 | 0.985468917 |
| 3.162278e-01 | 0.9947671 | 0.985887888 |
| 1.000000e+00 | 0.9952289 | 0.987123844 |
| 3.162278e+00 | 0.9950750 | 0.986706680 |
| 1.000000e+01 | 0.9950750 | 0.986706680 |
| 3.162278e+01 | 0.9950750 | 0.986706680 |
| 1.000000e+02 | 0.9950750 | 0.986706680 |

Accuracy was used to select the optimal model using the largest value.  
The final value used for the model was C = 1.

## Decision Trees:

```
# Evaluation Method
train_control = trainControl(method = "cv", number = 10)
# Fit the Model
tree1 <- train(type ~., data = wines, method = "rpart", trControl = train_control)
# Evaluate the Fit
tree1
```

CART

6497 samples  
12 predictor  
2 classes: 'red', 'white'

No pre-processing  
Resampling: Cross-Validated (10 fold)  
Summary of sample sizes: 5847, 5847, 5847, 5847, 5848, 5847, ...  
Resampling results across tuning parameters:

| cp         | Accuracy  | Kappa     |
|------------|-----------|-----------|
| 0.06253909 | 0.9448994 | 0.8427741 |
| 0.06754221 | 0.9344300 | 0.8125794 |
| 0.70043777 | 0.8370148 | 0.3857512 |


Accuracy was used to select the optimal model using the largest value.  
The final value used for the model was cp = 0.06253909.

Accuracy seems to be comparable in SVM and KNN, although not decision trees. I think the dataset may be too complex/large for a decision tree to handle. However, SVM and KNN seem to perform better. Although, due to the non-linear nature of this dataset, probably a KNN solution may be the best.

- e. Use the same already computed PCA again to show a scatter plot of the data and to visualize the labels for kNN, decision tree and SVM. Note that you do not need to recreate the PCA projection, you have already done this in 1b. Here, you just make a new visualization for each classifier using its labels for color (same points but change the color). Map the color results to the classifier, that is use the “predict” function to predict the class of your data, add it to your data frame and use it as a color. This is done for KNN in the tutorial, it should be similar for the others. Consider and explain the differences in how these classifiers performed.


### Using PCA:

```
# Plot and color the labels based on Wine type red (or) white
ggplot(data = rotated_data, aes(x = PC1, y = PC2, col = Color)) + geom_point(alpha = 0.3)
```



### Using KNN:


```
# Assign clusters as a new column
rotated_data$Clusters = as.factor(fit$cluster)
# Plot and color by labels
ggplot(data = rotated_data, aes(x = PC1, y = PC2 , col = Clusters)) + geom_point()
```



PCA seemed to depict the clusters into two uneven clusters, although we are aware that there is far more white wine samples than there is red wine in the dataset. The KNN clustering method provided an interesting distribution, although I prefer PCA as it seems to make more sense with the predicting of the red (or) white wine type. I think trying a different K value, perhaps smaller, may provide us with the better results. For the sake of testing it, I will try  $k = 2$  below, which seems to cluster the data more realistically:



```
fit2 <- kmeans(dummies, centers = 2, nstart = 25)
# Assign clusters as a new column
rotated_data$Clusters = as.factor(fit2$cluster)
# Plot and color by the labels
ggplot(data = rotated_data, aes(x = PC1, y = PC2, col = Clusters)) + geom_point()
```



## **PROBLEM – 2**

**In this question we will use the Sacramento data, which covers available housing in the region of that city. The variables include numerical information about the size of the housing and its price, as well as categorical information like zip code (there are a large but limited number in the area), and the type of unit (condo vs house (coded as residential))**

```
```{r}
data("Sacramento")
# Remove the Zipcode, and lon for simplicity
cpsacramento <- Sacramento %>% select(-c("latitude", "longitude", "zip"))
```
```

- a. Load the data from the tidyverse library with the data (“Sacramento”) command and you should have a variable Sacramento. Because we have categoricals, convert them to dummy variables.**

```
# Type is the largest variable to predict
dummy <- dummyVars(type ~., data = cpsacramento)
dummies <- as.data.frame(predict(dummy, newdata = cpsacramento))
head(dummies)
```

| Description: df [6 x 41] |               |             |                   |                 |                     |           |                      |                |                      |
|--------------------------|---------------|-------------|-------------------|-----------------|---------------------|-----------|----------------------|----------------|----------------------|
|                          | city.ANTELOPE | city.AUBURN | city.CAMERON_PARK | city.CARMICHAEL | city.CITRUS_HEIGHTS | city.COOL | city.DIAMOND_SPRINGS | city.EL_DORADO | city.EL_DORADO_HILLS |
| 1                        | 0             | 0           | 0                 | 0               | 0                   | 0         | 0                    | 0              | 0                    |
| 2                        | 0             | 0           | 0                 | 0               | 0                   | 0         | 0                    | 0              | 0                    |
| 3                        | 0             | 0           | 0                 | 0               | 0                   | 0         | 0                    | 0              | 0                    |
| 4                        | 0             | 0           | 0                 | 0               | 0                   | 0         | 0                    | 0              | 0                    |
| 5                        | 0             | 0           | 0                 | 0               | 0                   | 0         | 0                    | 0              | 0                    |
| 6                        | 0             | 0           | 0                 | 0               | 0                   | 0         | 0                    | 0              | 0                    |

6 rows | 1-10 of 41 columns

- b. With kNN, because of the high dimensionality, which might be a good choice for the distance function?

It's hard to tell without actually trying the different metrics, although using Minkowski is a fairly common in high dimensional data.

- c. Use kNN to classify this data with type as the label. Tune the choice of k plus the type of distance function. Report your results – what values for these parameters were tried, which were chosen, and how did they perform with accuracy?

```

```{r}
# Move the Type back to the dataset
sacramento_dummies <- dummies
sacramento_dummies$type <- Sacramento$type
library(kknn)
# Setup a tuneGrid with the tuning parameters
tuneGrid <- expand.grid(kmax = 3:7,          # Test a range of k values 3 to 7
                        kernel = c("rectangular", "cos"), # Regular and cosine-based distance funtions
                        distance = 1:3)      # Powers of Minkowski 1 to 3
# Tune and fit the model with 10-fold cross validation,
# Standardization, and our specialized tune grid
kknn_fit <- train(type ~.,
                  data = sacramento_dummies,
                  method = 'kknn',
                  trControl = ctrl,
                  preProcess = c('center', 'scale'),
                  tuneGrid = tuneGrid)
# Printing trained model provides report
kknn_fit
```

```

```

k-Nearest Neighbors

932 samples
41 predictor
3 classes: 'Condo', 'Multi_Family', 'Residential'

Pre-processing: centered (41), scaled (41)
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 838, 839, 839, 838, 838, 840, ...
Resampling results across tuning parameters:

  kmax kernel distance Accuracy Kappa
3    rectangular 1     0.9367516 0.3884981
3    rectangular 2     0.9367631 0.3763622
3    rectangular 3     0.9367631 0.3763622
3     cos        1     0.9356763 0.4495401
3     cos        2     0.9352556 0.4298051
3     cos        3     0.9324500 0.4205730
4    rectangular 1     0.9367516 0.3884981
4    rectangular 2     0.9367631 0.3763622
4    rectangular 3     0.9367631 0.3763622
4     cos        1     0.9388907 0.4619788
4     cos        2     0.9399660 0.4560239
4     cos        3     0.9399889 0.4391495
5    rectangular 1     0.9346011 0.3360441
5    rectangular 2     0.9303458 0.2866676
5    rectangular 3     0.9303458 0.2866676
5     cos        1     0.9388907 0.4619788
5     cos        2     0.9356992 0.3860980
5     cos        3     0.9378612 0.3943742
6    rectangular 1     0.9346011 0.3360441
6    rectangular 2     0.9303458 0.2866676
6    rectangular 3     0.9303458 0.2866676
6     cos        1     0.9367631 0.4273555
6     cos        2     0.9346240 0.3705842
6     cos        3     0.9378612 0.3943742
7    rectangular 1     0.9324620 0.3059469
7    rectangular 2     0.9335716 0.2803458
7    rectangular 3     0.9346354 0.2835647
7     cos        1     0.9367631 0.4273555
7     cos        2     0.9346240 0.3705842
7     cos        3     0.9389482 0.4006952

```

Accuracy was used to select the optimal model using the largest value.  
The final values used for the model were kmax = 4, distance = 3 and kernel = cos.

Here we tried different Minkowski distances from 1 to 3, and tested regular and cosine distance functions. K-max of 6, with a distance of 2 (Euclidean), and a cosine based distance function were the results. The cosine's distance function kappa value generally performed better, although accuracy was comparable on both.


### **PROBLEM – 3**

**In this problem we will continue with the wine quality data from Problem 1, but this time we will use clustering. Do not forget to remove the type variable before clustering because that would be cheating by using the label to perform clustering.**


```
# Copy wines dataset and remove type
cpwines <- wines
cpwines <- cpwines %>% select(-c("type"))
```

- a. Use k-means to cluster the data. Show your usage of silhouette and the elbow method to pick the best number of clusters. Make sure it is using multiple restarts.**

```
df <- cpwines
# Set seed
set.seed(123)
# Center scale allows us to standardize the data
preproc <- preProcess(df, method = c("center", "scale"))
# We have to call predict to fit our data based on preprocessing
predictors <- predict(preproc, df)
# Find the knee
fviz_nbclust(predictors, kmeans, method = "wss")
```



```
fviz_nbclust(predictors, kmeans, method = "silhouette")
```



The silhouette method suggests a  $k = 4$ , and the elbow method also suggests a  $k = 4$  may be reasonable.

```
# Fit the Data
fit <- kmeans(predictors, centers = 4, nstart = 25)
# Display the k means object information
fit
```

K-means clustering with 4 clusters of sizes 661, 1935, 1066, 2835

Cluster means:

|   | fixed.acidity | volatile.acidity | citric.acid | residual.sugar | chlorides  | free.sulfur.dioxide |
|---|---------------|------------------|-------------|----------------|------------|---------------------|
| 1 | 1.95495213    | 0.4048784        | 1.02423958  | -0.5697395     | 1.2673641  | -0.89472554         |
| 2 | -0.19278454   | -0.3500900       | 0.24718059  | 1.1525186      | -0.1081800 | 0.82522737          |
| 3 | 0.05693398    | 1.5882171        | -1.20795561 | -0.6076496     | 0.6266271  | -0.76925628         |
| 4 | -0.34563560   | -0.4526420       | 0.04668919  | -0.4253161     | -0.4572783 | -0.06538772         |

|   | total.sulfur.dioxide | density    | pH          | sulphates  | alcohol     | quality     |
|---|----------------------|------------|-------------|------------|-------------|-------------|
| 1 | -1.24231768          | 0.9050474  | -0.11205327 | 1.3614349  | 0.04704102  | 0.03647183  |
| 2 | 0.95020983           | 0.7315214  | -0.38723928 | -0.2693766 | -0.80429691 | -0.29268305 |
| 3 | -1.03454506          | 0.4579506  | 0.88616595  | 0.3628512  | -0.28754458 | -0.57836111 |
| 4 | 0.03010265           | -0.8825064 | -0.04277873 | -0.2700050 | 0.64611743  | 0.40873607  |

Clustering vector:

```
[1] 3 3 3 1 3 3 3 3 3 3 3 3 1 1 1 1 1 3 1 1 3 1 3 3 3 3 1 3 3 3 3 3 3 3 1 3 3 3 3 1 3
[45] 3 3 3 1 3 3 3 3 3 1 3 3 1 3 3 3 3 1 3 3 3 3 3 3 3 3 3 3 1 1 3 3 3 3 1 3 1 1 3 1 3
[89] 1 3 3 1 1 3 3 3 3 3 3 3 3 3 3 3 3 1 3 1 1 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 1 3
[133] 3 3 3 3 3 3 3 3 3 3 4 3 4 1 3 1 3 1 1 1 3
[177] 3 3 3 3 1 3 3 3 1 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 1 3 3 3 1 1 3 3 1 3 3 3 3
[221] 1 3 3 3 3 3 1 3
[265] 1 1 3 1 3 1 3 1 3 1 3 3 3 3 1 1 1 1 1 3 3 3 3 3 3 3 3 3 1 3 3 1 1 1 3 3 3 3 3 3 3
[309] 1 3 1 3 1 3 3 3 3 3 1 3 1 3 3 1 1 1 1 1 1 1 1 3 3 3 3 1 3 3 1 1 1 1 1 1 3 3 1 3 3
[353] 3 1 4 3 1 1 1 1 3 1 1 1 1 1 1 1 1 1 3 1 3 1 1 1 1 1 1 1 1 1 1 3 3 3 3 3 3 1 3 1 1 1
[397] 3 1 1 3 3 4 1 1 3 1 1 1 1 1 1 1 1 3 1 3 1 3 1 3 3 3 3 3 3 3 3 3 3 3 1 1 3 1 1 1 3 3
[441] 1 1 1 1 3 3 1 1 3 1 1 1 3 1 4 1 1 3 1 1 1 3 1 1 1 1 1 1 1 1 3 3 1 1 1 1 1 3 1 3 1 1
[485] 1 1 1 1 1 1 3 1 1 3 4 1 3 3 1 3 3 1
[529] 1 1 1 1 1 1 1 1 1 3 1 1 3 1 3 1 1 3 1 1 3 1 1 3 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 3 3
[573] 1 1 1 3 1 1 1 1 1 1 1 1 1 3 1 3 4 1 1 4 1 1 3 1 1 3 1 3 1 3 3 1 1 1 3 1 1 3 1 1 1
[617] 1 1 1 1 3 3 3 3 3 3 3 3 3 3 3 3 1 3 3 3 3 3 3 3 1 1 1 1 1 3 3 3 1 4 1 3 1 1 3 1 1 3
[661] 3 3 3 1 1 3 1 1 1 1 3 3 3 3 1 1 1 3 3 1 1 3 3 3 3 3 3 3 3 3 1 3 3 1 3 3 3 3 3 1 3 3
[705] 3 3 3 3 3 1 1 3 3 3 3 3 3 3 3 3 3 3 1 3
[749] 3 3 3 3 3 1 3
[793] 3 3 1 1 1 1 1 3 3 3 3 3 3 1 1 1 3 3 3 1 1 3 1 1 1 1 3 3 3 3 3 3 3 3 4 3 3 3 3 1 3 3
[837] 4 4 1 3 1 3 1 3 1 3 3 3 3 3 1 1 1 1 3 1 1 1 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 1 1 3 3 3
[881] 3 3 1 3 3 3 3 1 3 1 3 3 1 3 3 3 1 3 1 3 1 3 3 3 3 3 3 3 3 3 4 1 1 1 4 1 3 3 1 1 3 3
[925] 1 1 1 3 1 1 3 3 3 3 1 1 1 4 3 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 3 1 1 1 3 3 1 3 1 1 1 3
[969] 1 3 1 1 1 1 3 3 3 4 1 1 3 4 1 1 3 1 3 1 3 1 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
[ reached getOption("max.print") -- omitted 5497 entries ]
```

Within cluster sum of squares by cluster:  
[1] 8421.783 13176.504 7011.810 17841.974  
(Between\_SS / total\_SS = 40.4 %)

Available components:

```
[1] "cluster" "centers" "totss" "withinss" "tot.withinss" "betweenss"
[7] "size" "iter" "ifault"
```

- b. Use hierarchical agglomerative clustering (HAC) to cluster the data. Try at least 2 distance functions and at least 2 linkage functions (cluster distance functions), for a total of 4 parameter combinations. For each parameter combination, perform the clustering.

### **Euclidean and complete linkage:**

```
dist_mat <- dist(predictors, method = 'euclidean')
# Determine assembly/agglomeration method and run hclust
hfit1 <- hclust(dist_mat, method = 'complete')
hfit1
```

```
Call:
hclust(d = dist_mat, method = "complete")

Cluster method      : complete
Distance            : euclidean
Number of objects   : 6497
```

### **Euclidean and average linkage:**

```
dist_mat <- dist(predictors, method = 'euclidean')
# Determine assembly/agglomeration method and run hclust
hfit2 <- hclust(dist_mat, method = 'average')
hfit2
```

```
Call:
hclust(d = dist_mat, method = "average")

Cluster method      : average
Distance            : euclidean
Number of objects   : 6497
```

### **Manhattan and complete linkage:**

```
dist_mat <- dist(predictors, method = 'manhattan')
# Determine assembly/agglomeration method and run hclust (average uses mean)
hfit3 <- hclust(dist_mat, method = 'complete')
hfit3
```

```
Call:
hclust(d = dist_mat, method = "complete")

Cluster method      : complete
Distance            : manhattan
Number of objects   : 6497
```

### **Manhattan and average linkage:**

```
dist_mat <- dist(predictors, method = 'manhattan')
# Determine assembly/agglomeration method and run hclust (average uses mean)
hfit4 <- hclust(dist_mat, method = 'average')
hfit4
```

```
Call:
hclust(d = dist_mat, method = "average")

Cluster method      : average
Distance            : manhattan
Number of objects   : 6497
```

```
# Build the new model
h1 <- cutree(hfit1, k=4)
h2 <- cutree(hfit2, k=4)
h3 <- cutree(hfit3, k=4)
h4 <- cutree(hfit4, k=4)
```

- c. Compare the k-means and HAC clusterings by creating a crosstabulation between their labels.

```
#Redefining the fit (I think I missed it previously)
fit <- kmeans(predictors, centers = 4, nstart = 25)
#Create a dataframe for the results
result1 <- data.frame(WineType = wines$type, HAC1 = h1, Kmeans = fit$cluster)
result2 <- data.frame(WineType = wines$type, HAC2 = h2, Kmeans = fit$cluster)
result3 <- data.frame(WineType = wines$type, HAC3 = h3, Kmeans = fit$cluster)
result4 <- data.frame(WineType = wines$type, HAC4 = h4, Kmeans = fit$cluster)
```

```
#Crosstab for HAC
result1 %>% group_by(HAC1) %>% select(HAC1, WineType) %>% table()
```

|      | WineType |       |
|------|----------|-------|
| HAC1 | red      | white |
| 1    | 1597     | 4896  |
| 2    | 2        | 0     |
| 3    | 0        | 1     |
| 4    | 0        | 1     |

```
result2 %>% group_by(HAC2) %>% select(HAC2, WineType) %>% table()
```

|      | WineType |       |
|------|----------|-------|
| HAC2 | red      | white |
| 1    | 1575     | 4895  |
| 2    | 24       | 1     |
| 3    | 0        | 1     |
| 4    | 0        | 1     |

```
result3 %>% group_by(HAC3) %>% select(HAC3, WineType) %>% table()
```

|      | WineType |       |
|------|----------|-------|
| HAC3 | red      | white |
| 1    | 1595     | 4889  |
| 2    | 4        | 0     |
| 3    | 0        | 8     |
| 4    | 0        | 1     |

```
result4 %>% group_by(HAC4) %>% select(HAC4, WineType) %>% table()
```

|      | WineType |       |
|------|----------|-------|
| HAC4 | red      | white |
| 1    | 1576     | 4896  |
| 2    | 23       | 0     |
| 3    | 0        | 1     |
| 4    | 0        | 1     |


```
#Crosstab for K Means
result <- data.frame(Type = wines$type, Kmeans = fit$cluster)
result %>% group_by(Kmeans) %>% select(Kmeans, Type) %>% table()
```

|        | Type |       |
|--------|------|-------|
| Kmeans | red  | white |
| 1      | 3    | 1932  |
| 2      | 609  | 52    |
| 3      | 927  | 139   |
| 4      | 60   | 2775  |

- d. For comparison – use PCA to visualize the data in a scatterplot. Create 3 separate plots: use the color of the points to show (1) the type label, (2) the k-means cluster labels and (3) the HAC cluster labels.


### PCA:

```
#Recreating the PCA scatter plot
#Create Dummies
dummy <- dummyVars(type ~ ., data = wines)
dummies <- as.data.frame(predict(dummy, newdata = wines))
set.seed(123)
#Calculate PCA
pca = prcomp(dummies)
#Save as data frame
rotated_data = as.data.frame(pca$x)
#Add original label 'type' as a reference
rotated_data$Color <- wines$type
#Plot and color the labels based on wine type red (or) white
ggplot(data = rotated_data, aes(x = PC1, y = PC2, col = Color)) + geom_point(alpha = 0.3)
```




### H3:

```
rotated_data$Clusters = as.factor(h3)
# Plot and color by labels
ggplot(data = rotated_data, aes(x = PC1, y = PC2, col = Clusters)) + geom_point()
```



### **K-means:**

```
rotated_data$Clusters = as.factor(fit$cluster)
# Plot and color by labels
ggplot(data = rotated_data, aes(x = PC1, y = PC2, col = Clusters)) + geom_point()
```



- e. **Consider the results of C and D and explain the differences between the clustering results in terms of how the algorithms work.**

K-means uses a pre-specified K value, while HAC doesn't. We can see that the clustering method seems to be less random in the K-means method. Since HAC is arranged more like a 'tree', it seemed to struggling with clustering the data. K-means divided the data into non-overlapping clusters, which provided a more reasonable look at the dataset. Regardless, I still find PCA to be more sensible visually in this



case, even though I don't think it performed as well as k-means did in this case as it didn't handle the complexity of the dataset as well.

#### **PROBLEM – 4**

**Back to the Starwars data from a previous assignment! Remember that the variable that lists the actual names and the variables that are actually lists will be a problem, so remove them (name, films, vehicles, starships). Make sure to double check the types of the variables, i.e., that they are numerical or factors as you expect.**


```
```{r}
data("starwars")
#Copy Starwars
cpstarwars <- starwars
#Remove some columns
cpstarwars <- cpstarwars %>% select(-c("name", "vehicles", "starships", "films"))
#Remove NAs
cpstarwars <- na.omit(cpstarwars)
summary(cpstarwars)
```
```

|                  |                  |                  |                  |                  |                |
|------------------|------------------|------------------|------------------|------------------|----------------|
| height           | mass             | hair_color       | skin_color       | eye_color        | birth_year     |
| Min. : 88        | Min. : 20.00     | Length:29        | Length:29        | Length:29        | Min. : 8.00    |
| 1st Qu.:170      | 1st Qu.: 75.00   | Class :character | Class :character | Class :character | 1st Qu.: 31.00 |
| Median :180      | Median : 79.00   | Mode :character  | Mode :character  | Mode :character  | Median : 46.00 |
| Mean :178        | Mean : 77.77     |                  |                  |                  | Mean : 51.29   |
| 3rd Qu.:188      | 3rd Qu.: 83.00   |                  |                  |                  | 3rd Qu.: 57.00 |
| Max. :228        | Max. :136.00     |                  |                  |                  | Max. :200.00   |
| sex              | gender           | homeworld        | species          |                  |                |
| Length:29        | Length:29        | Length:29        | Length:29        |                  |                |
| Class :character | Class :character | Class :character | Class :character |                  |                |
| Mode :character  | Mode :character  | Mode :character  | Mode :character  |                  |                |

- a. Use hierarchical agglomerative clustering to cluster the Starwars data. This time we can leave the categorical variables in place, because we will use the gower metric from daisy in the cluster library to get the distances. Use average linkage. Determine the best number of clusters.**

I had some trouble trying to run this without converting the categoricals to dummies, so I decided to proceed using the dummy variables.

```
library(cluster)
#Pass dataframe directly with metric = gower
dist_mat <- daisy(dummies, metric = "gower")
#Center scale allows us to standardize the data
preproc <- predict(preproc, cpstarwars)
#Silhouette score comparison to find K
fviz_nbclust(predictors, FUN = hcut, method = "silhouette")
```




### Clusters:

```
#Determine the assembly/agglomeration method and run hclust
hfit <- hclust(dist_mat, method = 'average')
#Build the new model
h2 <- cutree(hfit, k=2)
summary(h2)
```

```
Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 1         1         1         1         1         2
```

- b. Produce the dendrogram for (a). How might an anomaly show up in a dendrogram? Do you see a Starwars character who does not seem to fit in easily? What is the advantage of considering anomalies this way as opposed to looking for unusual values relative to the mean and standard deviations, as we considered earlier in the course? Disadvantages?

```
hfit <- hclust(dist_mat, method = 'average')
plot(hfit)
```



Own branch with no relationship to other members of the dataset. Based on this dataset, 21, 9, and 18 may have anomalous features. The advantages of this dendrogram is the ability to view anomalies directly without reviewing the actual dataset in a tabulation. However, we're only able to view the anomalies here based on height, which may be lacking in regards to other features like eye color (or) mass.

- c. Use dummy variables to make this data fully numeric and then use k-means to cluster. Choose the best number of clusters.


```
dummy <- dummyVars(gender ~ ., data = cpstarwars)
dummies <- as.data.frame(predict(dummy, newdata = cpstarwars))
head(dummies)
```

Description: df [6 x 66]


|   | height<br><dbl> | mass<br><dbl> | hair_colorauburn, white<br><dbl> | hair_colorblack<br><dbl> | hair_colorblond<br><dbl> | hair_colorbrown<br><dbl> |
|---|-----------------|---------------|----------------------------------|--------------------------|--------------------------|--------------------------|
| 1 | 172             | 77            | 0                                | 0                        | 1                        | 0                        |
| 2 | 202             | 136           | 0                                | 0                        | 0                        | 0                        |
| 3 | 150             | 49            | 0                                | 0                        | 0                        | 1                        |
| 4 | 178             | 120           | 0                                | 0                        | 0                        | 0                        |
| 5 | 165             | 75            | 0                                | 0                        | 0                        | 1                        |
| 6 | 183             | 84            | 0                                | 1                        | 0                        | 0                        |

6 rows | 1-7 of 66 columns

```
#Create a predictors file using dummies
predictors <- dummies
#Set seed
set.seed(123)
#Center scale allows us to standardize the data
preproc <- preProcess(predictors, method = c("center", "scale"))
#We have to call predict to fit our data based on preprocessing
predictors <- predict(preproc, predictors)
#Find the knee
fviz_nbclust(predictors, kmeans, method = "wss")
```



```
fviz_nbclust(predictors, kmeans, method = "silhouette")
```



K=2 will be the distance we are using here,

```
#Fit the data
fit <- kmeans(predictors, centers = 2, nstart = 25)
#Display the kmeans object information
fit
```

K-means clustering with 2 clusters of sizes 9, 20

```
Cluster means:
  height      mass hair_colorauburn, white hair_colorblack hair_colorblond hair_colorbrown hair_colorbrown, grey
1 0.6186563 0.6114774 -0.1856953 -0.5018706 -0.2674319 -0.2991273 -0.1856953
2 -0.2784039 -0.2751649 0.0835629 0.2258418 0.1203443 0.1346079 0.0835629

  hair_colorgrey hair_colornone hair_colorwhite skin_colorblue skin_colorbrown skin_colorbrown mottle skin_colordark
1 0.4126563 0.7568054 0.16343058 -0.1856953 -0.1856953 0.4126563 -0.2674319
2 -0.1856953 -0.3405624 -0.07354376 0.0835629 0.0835629 -0.1856953 0.1203443

  skin_colorfair skin_colorgreen skin_colorlight skin_colororange skin_colorpale skin_colorred skin_colortan
1 0.5542653 0.4126563 -0.5018706 0.5942930 0.5942930 0.4126563 -0.1856953
2 0.2494194 -0.1856953 0.2258418 -0.2674319 -0.2674319 -0.1856953 0.0835629

  skin_colorunknown skin_colorwhite skin_coloryellow eye_colorblack eye_colorblue eye_colorblue-gray eye_colorbrown
1 0.4126563 0.4126563 -0.2674319 0.4126563 -0.3622024 -0.1856953 -0.1228583
2 -0.1856953 -0.1856953 0.1203443 -0.1856953 0.1628911 0.0835629 0.3077662

  eye_colorhazel eye_colororange eye_colorred eye_coloryellow birth_year sexfemale sexmale homeworldAlderaan
1 -0.2674319 0.5942930 0.4126563 0.8734288 0.5428396 -0.5018706 0.5018706 -0.1856953
2 0.1203443 -0.2674319 -0.1856953 -0.3930428 -0.2442778 0.2258418 -0.2258418 0.0835629

  homeworldBespin homeworldCerea homeworldConcord Dawn homeworldCorellia homeworldDathomir homeworldDorin
1 -0.1856953 0.4126563 -0.1856953 -0.2674319 0.4126563 0.4126563 0.4126563
2 0.0835629 -0.1856953 0.0835629 0.1203443 -0.1856953 -0.1856953 -0.1856953

  homeworldEndor homeworldHarun Kai homeworldHoth homeworldHuyyik homeworldMirial homeworldMon Cala
1 -0.1856953 -0.1856953 -0.1856953 0.4126563 -0.2674319 0.4126563
2 0.0835629 0.0835629 0.0835629 -0.1856953 0.1203443 -0.1856953

  homeworldNaboo homeworldRyloth homeworldSerrenno homeworldSocorro homeworldStewjon homeworldTatooine
1 0.3832233 -0.1856953 -0.1856953 -0.1856953 -0.1856953 -0.2323475
2 -0.1724505 0.0835629 0.0835629 0.0835629 0.0835629 0.1045564

  homeworldTrandosha speciesCerean speciesEwok speciesGungan speciesHuman speciesKel Dor speciesMirialan
1 0.4126563 0.4126563 -0.1856953 0.4126563 -0.8069344 0.4126563 -0.2674319
2 -0.1856953 -0.1856953 0.0835629 -0.1856953 0.3631205 -0.1856953 0.1203443

  speciesMon Calamari speciesTrandoshan speciesWilek speciesWoodree speciesZabrak
1 0.4126563 0.4126563 -0.1856953 0.4126563 0.4126563
2 -0.1856953 -0.1856953 0.0835629 -0.1856953 -0.1856953
```

```
Clustering vector:
1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29
2 1 2 2 2 2 2 2 2 1 2 2 1 2 1 2 2 1 2 1 1 2 2 1 1 2 2 2 2 2
```

```
Within cluster sum of squares by cluster:
[1] "18.4747 985.1128"
(between_SS / total_SS = 7.8 %)
```

Available components:

```
[1] "cluster" "centers" "totss" "withinss" "tot.withinss" "betweenss" "size"
[8] "iter"
```

**d. Compare the HAC and k-means clusterings with a crosstabulation.**

```
```{r}
#Create a dataframe for results
result <- data.frame(Gender = cpstarwars$gender, HAC2 = h2, Kmeans = fit$cluster)
#Create a cross tab for HAC
result %>% group_by(HAC2) %>% select(HAC2, Gender) %>% table()
```
```